



Content

- 01** Problem
- 02** Solution
- 03** Dataset
- 04** Model Approach
- 05** Model Deployment
- 06** Architecture Plan
- 07** Results & Evaluation
- 08** Conclusion



Problem



The main objective is to maximize the bank's revenue by minimizing the **unnecessary retention programs** and by **minimizing the marketing and acquisition costs** for new customers by improving the churn model prediction

Solution




Data Preparation

Collect and Clean Data: Assemble historical customer data and clean for quality

Feature Engineering: Develop features indicative of churn, such as account activity patterns

Data Labeling: Define and apply labels for 'churned' or 'retained' status



Model Deployment

Algorithm Selection: Choose predictive algorithms suited to the data and problem

Model Training: Train the model on historical data, validating with a separate test and validation set

Performance Evaluation: Assess the model using metrics like F1 score and ROC-AUC



Implementation & Monitoring

Deployment: Integrate the model into the bank's operational systems

Monitoring: Continuously monitor the model's performance and update as needed

Refinement: Use model insights to refine and target retention strategies effectively

Bank Customer Churn Dataset




Target variable: churn

1 if the client has left the bank during some period or 0 if he/she has not



Customer ID: a unique identifier for each customer



Numerical variables

Credit score	From 350 to 850
Age	From 18 to 92
Tenure	From 0 to 10
Balance	Max: 251k
Product numbers	From 1 to 4
Estimated salary	Max: 200k



Categorical variables

Country	Spain, France, Germany
Gender	55% male, 45% female
Credit card	70,55% have a credit card
Active membership	51,51% are active members

Most of the variables have bell-shaped distribution, except for tenure and estimated salary, which are more uniform-like

Data Cleaning



There are no **missing values**, so no imputation methods are required.

Customer_id is irrelevant for churn prediction, indicating it will be **removed** before model training.

The correlation matrix indicates a lack of significant correlations between features, negating the need for feature selection or dimensionality reduction.


customer_id	1	0.01	0.01	-0.01	-0.01	0.02	-0.01	0	0.02	-0.01
credit_score	0.01	1	-0	0	0.01	0.01	-0.01	0.03	-0	-0.03
age	0.01	-0	1	-0.01	0.03	-0.03	-0.01	0.09	-0.01	0.29
tenure	-0.01	0	-0.01	1	-0.01	0.01	0.02	-0.03	0.01	-0.01
balance	-0.01	0.01	0.03	-0.01	1	-0.3	-0.01	-0.01	0.01	0.12
products_number	0.02	0.01	-0.03	0.01	-0.3	1	0	0.01	0.01	-0.05
credit_card	-0.01	-0.01	-0.01	0.02	-0.01	0	1	-0.01	-0.01	-0.01
active_member	0	0.03	0.09	-0.03	-0.01	0.01	-0.01	1	-0.01	-0.16
estimated_salary	0.02	-0	-0.01	0.01	0.01	0.01	-0.01	-0.01	1	0.01
churn	-0.01	-0.03	0.29	-0.01	0.12	-0.05	-0.01	-0.16	0.01	1
	customer_id	credit_score	age	tenure	balance	products_number	credit_card	active_member	estimated_salary	churn

Model Approach

For the ML model a **XGBoost** model was implemented

XGBoost is an **advanced implementation of gradient boosted decision trees** designed for speed and performance

It utilizes **regularized boosting and advanced tree learning** algorithms to improve accuracy and prevent overfitting



Why XGBoost?

Superior Results

Prevents overfitting, ensuring reliable churn predictions


Speed and Scale

Efficiently processes large customer datasets

Insightful Analytics

Offers clear feature importance for targeted strategies

Model Deployment



Load the data into a python notebook

Pre-process the data

Implement the XGBoost model

Upload the Notebook into AWS SageMaker



Train the model in AWS SageMaker

Evaluate and check the metrics of the trained model

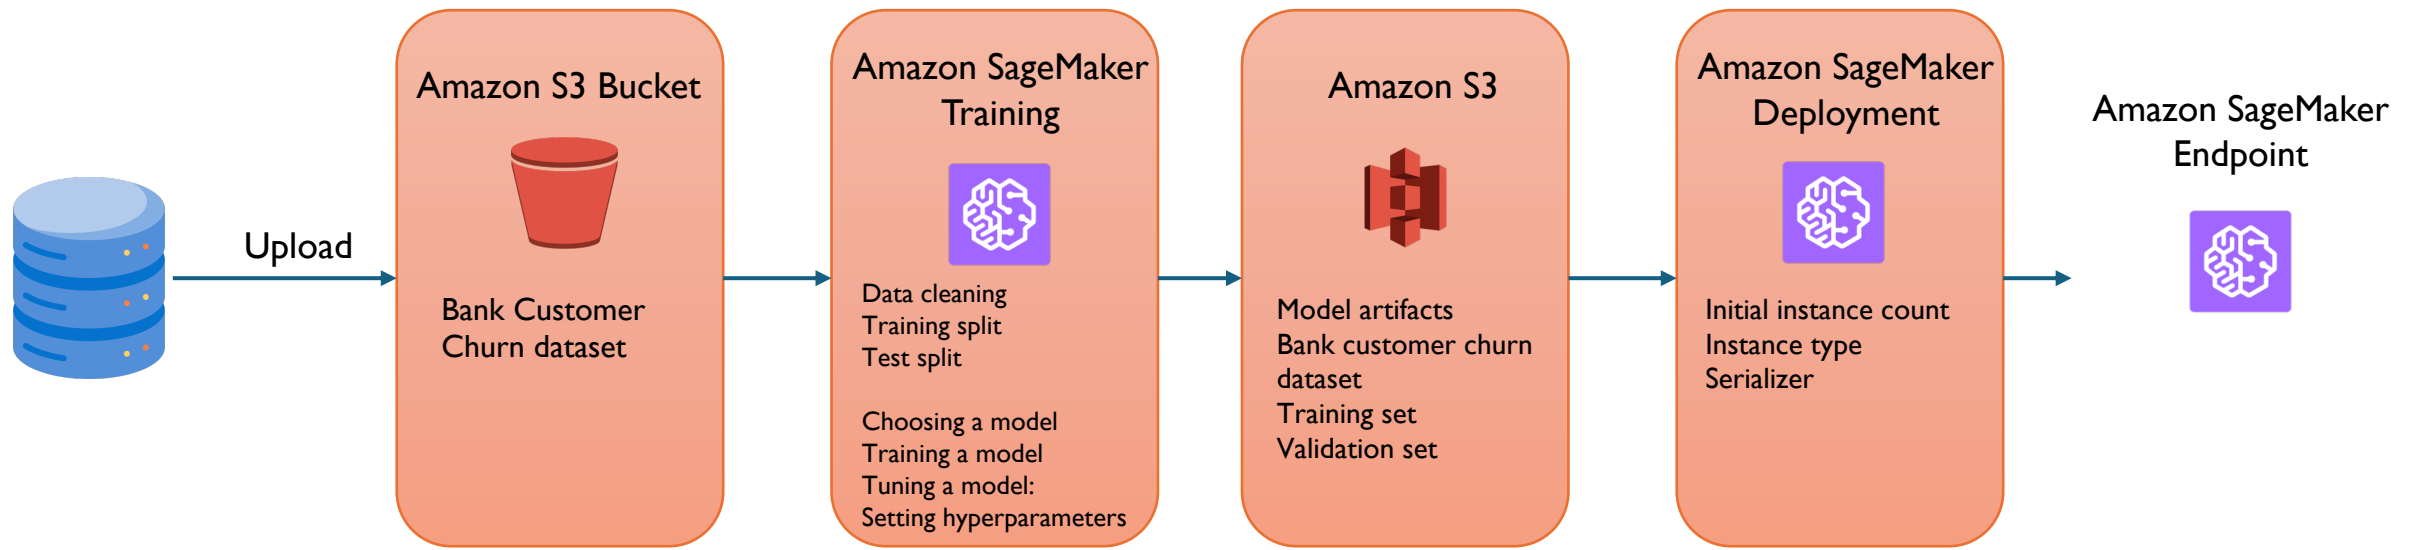
Configure and launch a hyperparameter tuning job

Deployment

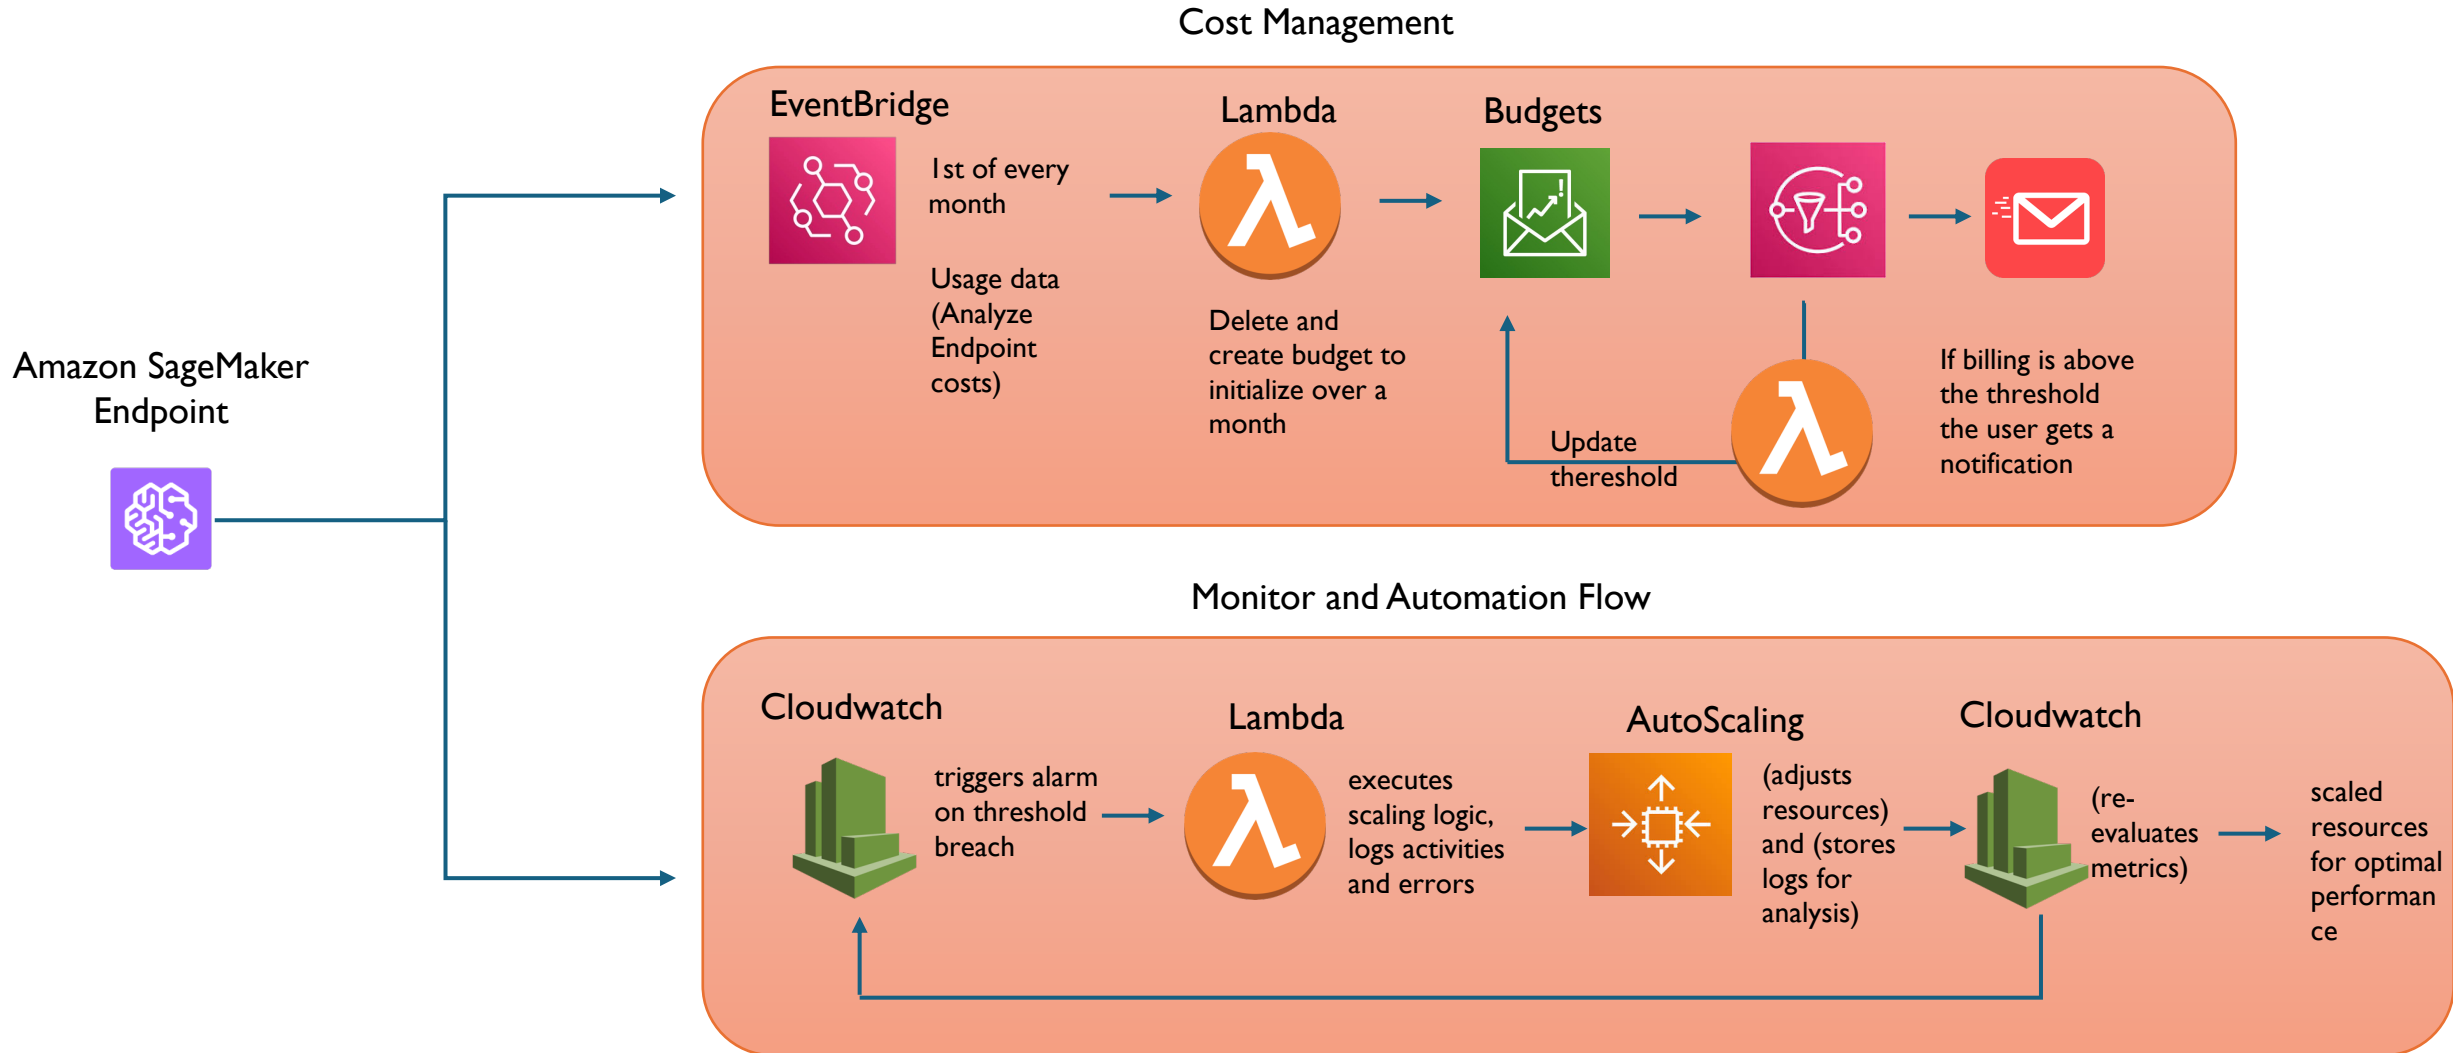
Find the optimal hyperparameter values that maximize the area under the ROC curve (AUC) on the validation dataset

Evaluate the results of the tuned model and deploy it

Architecture Diagram



Next Steps



Results & Evaluation

Train AUC: 0.8983
Validation AUC: 0.8518



Define Key & fixed Hyperparameters
And the hyperparameter range



Train AUC: 0.9084
Validation AUC: 0.8537

Accuracy: High accuracy indicated that the model effectively classifies churn and non-churn cases

84%

Precision: High precision indicates the model is reliable in its churn predictions

72%

Recall: There's potential to improve the model's detection of actual churn cases

41%

F1 Score: Suggests enhancing recall is necessary for better churn identification

52%



Insights:

- **Improved accuracy** through effective hyperparameter exploration
- Majority of predicted probabilities are **below the 0.5 threshold**, suggesting low churn likelihood
- The model shows **robust overall performance** but has room for improvement in sensitivity to churn

Conclusion

Churn Model Impact


Implemented a predictive churn model using XGBoost that achieved solid rate of effectively identifying customers at risk of leaving the bank

Data Utilization

Leveraged comprehensive data preparation and model training to highlight the most influential factors affecting customer churn

Efficiency Gain

Streamlined retention programs and reduced unnecessary costs, resulting in a more targeted and financially efficient marketing strategy



Next Steps

Enhanced Automation

Deploy a DynamoDB or RDS in combination with a lambda function to make the database modification more accessible

Model Iteration

Experiment with additional hyperparameter tuning, ensemble methods, or alternative algorithms

Customer Engagement

Integrate the model's insights with customer engagement strategies to proactively address factors contributing to churn

Thank you!